

AI assist (optional)

OpenTrack can use an AI model to help with issues — starting with **smart triage** (suggesting a severity, priority, category, and tags for a new issue from what you typed). It's **off by default** and only ever runs when you turn it on and point it at a provider. Whatever it suggests is just a suggestion — a person always accepts or changes it.

You are **not** locked into one AI company. OpenTrack talks to two kinds of provider, so you can pick what fits your budget and privacy needs.

Which AI can I use?

You want...	Provider setting	Notes
Anthropic Claude (what the author uses)	<code>anthropic</code>	Cloud. Needs an Anthropic API key.
OpenAI (ChatGPT models, e.g. GPT-4o mini)	<code>openai</code>	Cloud. Needs an OpenAI API key.
Azure OpenAI	<code>openai + BaseUrl</code>	Cloud, your Azure resource.
Groq / OpenRouter / other hosted	<code>openai + BaseUrl</code>	Cloud. Any OpenAI-compatible endpoint.
Ollama (free, runs on your own PC)	<code>openai + BaseUrl</code>	Local — no key, no data leaves your machine.
LM Studio (free, runs on your own PC)	<code>openai + BaseUrl</code>	Local — same privacy benefit.

The `openai` setting means "any service that speaks OpenAI's Chat Completions format." A huge number of tools — including the free local ones — do, which is why one setting covers so many options.

Important: who pays?

The cloud providers bill an **API account** at that provider — this is **separate** from any monthly chat subscription. In particular, an Anthropic API key is billed to your **Anthropic API account**, which is **not** the same as a Claude Pro/Max subscription; turning on AI here does **not** draw from a Claude.ai plan. The **local** options (Ollama, LM Studio) are **free** and run on your own computer — no account, no per-use charge.

Cloud costs for these features are small — a triage suggestion is typically a fraction of a cent — but always check the provider's current pricing, and set a spend limit where the provider offers one.

Step-by-step: get an Anthropic (Claude) API key

- 1 Open a browser and go to `<https://console.anthropic.com>`. This is the

Anthropic Console (the developer/API site) — **not** claude.ai, the chat site. They are separate accounts even if you use the same email.

- 1 **Sign in**, or click **Sign up** and create an account (email + password, or

"Continue with Google"). Verify your email if asked.

- 1 Once you're in the Console, look at the **left sidebar** (or the gear/■

Settings menu) and click **API Keys**. (Direct link: <<https://console.anthropic.com/settings/keys>>.)

- 1 Click **Create Key** (sometimes labeled "Create API Key").
- 2 Give it a name you'll recognize later, like `OpenTrack`, and click **Create**

/ **Add**.

- 1 The key is shown **once**. It looks like `sk-ant-api03-xxxxxxx...`. Click

Copy and paste it somewhere safe **right now** — you can't view it again later (only delete it and make a new one).

- 1 **Add credit so the key works**. A brand-new API account usually has a \$0

balance and calls will fail until you fund it. In the sidebar go to **Billing** → **Plans / Buy credits**, add a payment method, and buy a small amount (even \$5 is plenty for triage). While you're there, set a **monthly spend limit** so it can never surprise you.

- 1 Keep that `sk-ant-...` key handy for the "Turn it on" step below.

Step-by-step: get an OpenAI API key

- 1 Go to <<https://platform.openai.com>> — the **OpenAI developer platform**,

not chatgpt.com. Different account from a ChatGPT Plus subscription.

- 1 **Sign in** or **Sign up** and verify your email (and phone, if asked).
- 2 Click your **profile icon (top-right)** → **View API keys**, or go straight to

<<https://platform.openai.com/api-keys>>.

- 1 Click **Create new secret key**, name it `OpenTrack`, and click **Create**.
- 2 Copy the key (it starts with `sk-...`) **now** — it's shown only once.
- 3 Go to **Settings** → **Billing**, add a payment method, and add a little credit.

Set a **usage limit** while you're there.

- 1 Pick a cheap, capable model name for the config below — `gpt-4o-mini` is a good default.

Free & private option: run it locally with Ollama (no key at all)

If you'd rather **not** send issue text to any cloud, run a model on your own machine. Nothing leaves your computer.

- 1 Download and install **Ollama** from [<https://ollama.com>](https://ollama.com) (Windows, Mac, Linux).

- 1 Open a terminal and pull a model, e.g.:

```
ollama pull llama3.1
```

Ollama then serves an OpenAI-compatible API at `http://localhost:11434/v1`.

- 1 Use the "Local (Ollama)" settings below — **no API key needed**.

(LM Studio works the same way; it serves on `http://localhost:1234/v1`.)

Turn it on

Put the settings in configuration — `appsettings.json`, environment variables, or .NET user-secrets.

Prefer environment variables or user-secrets for the key so it isn't committed to source control. As environment variables, replace each `:` with `__` (double underscore), e.g.

```
OpenTrack__Ai__Enabled=true.
```

Claude (Anthropic):

```
OpenTrack:Ai:Enabled = true
OpenTrack:Ai:Provider = anthropic
OpenTrack:Ai:ApiKey   = sk-ant-...           (your Anthropic key)
OpenTrack:Ai:Model    = claude-haiku-4-5-20251001 (fast & inexpensive)
```

OpenAI (cloud):

```
OpenTrack:Ai:Enabled = true
OpenTrack:Ai:Provider = openai
OpenTrack:Ai:ApiKey   = sk-...               (your OpenAI key)
OpenTrack:Ai:Model    = gpt-4o-mini
```

Local, free & private (Ollama):

```
OpenTrack:Ai:Enabled = true
OpenTrack:Ai:Provider = openai
OpenTrack:Ai:BaseUrl  = http://localhost:11434/v1
OpenTrack:Ai:Model    = llama3.1
                                                                (no ApiKey needed for local)
```

Azure OpenAI / Groq / OpenRouter / other hosted: use `Provider = openai`, set `BaseUrl` to that service's endpoint, `ApiKey` to its key, and `Model` to a model it offers.

Restart OpenTrack. On the **New issue** page you'll now see a **■ Suggest with AI** button that fills in severity/priority/category and proposes tags.

Privacy & safety

- **Opt-in.** With AI off (the default) or unconfigured, OpenTrack behaves exactly

as before and makes no AI calls at all.

- **Cloud providers see the text.** With a cloud provider on, a suggestion sends

that issue's summary/description to that provider. Don't enable a cloud provider for projects whose contents can't leave your environment — use a **local** engine (Ollama/LM Studio) instead, which keeps everything on your machine.

- **Server-side only.** The key is read from server configuration and is never

sent to the browser or stored in the database.

- **Best-effort.** Every AI result is a suggestion with a human in the loop; if a

call fails, issue creation is completely unaffected.